

# Kmeans Clustering

Clustering is an unsupervised learning method as we don't have labels for the dataset. We try to find similarity by forming cluster in an n-dimensional space. Points which are "closer" to each other are considered more similar than the ones which are further apart. This "closeness" can be measure in a number of ways like Euclidian distance, cosine similarity etc.

Kmeans is one of the most popular clustering algorithm. Our objective is to find k points in the n-dimensional space such these points sit at the centre of the k dense clusters formed by plotting the dataset in n-dimensional space. The basic implementation of the algorithm is as follows

- 1) Select k unique random points from the dataset and call them initial centroids
- 2) For each point in the dataset find the closest centroid and assign that point to the cluster of that centroid
- 3) For each cluster, calculate the mean of all the points. Mean values will give new centroids
- 4) Repeat steps 2 and 3 until you get a new set of centroids which is very close to the previous one.

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from PIL import Image
from sklearn.datasets.samples_generator import make_blobs
```

```
In [2]: def initialize_centroids(X,k):
        """
        randomly picks k unique points from X as the initial centroids

        inputs--
        X: input matrix--> np.ndarray
        k: number of centroids

        outputs--
        array of centroid vectors
        """
        return X[np.random.choice(range(X.shape[0]), k, replace=False)]
```

```
In [3]: def distance_bw_centroids(old:np.ndarray,new:np.ndarray):
        """returns euclidian distance between old and new centroids"""
        return np.linalg.norm(old.ravel()-new.ravel())
```

```
In [4]: def assign_buckets(X:np.ndarray, centroids:np.ndarray):
        """
        returns a np.array of bucket indices
        for a given set of points and centroids
        """
        dist_matrix = np.stack([np.linalg.norm(X-c, axis=1) for c in centroids])
        return np.argmin(dist_matrix, axis=0)
```

```
In [5]: def kmeans(X:np.ndarray, k:int, centroids=None, tolerance=1e-2):
        """
        Inputs :
            X:Dataset array
            k:number of clusters
            centroids: None for random initialization, 'kmeans++' for smart
            initialization
            tolerance: max tolerable distance between new and old centroids
        Outputs :
            array of centroids (k,X.shape[1])
            array of cluster indices corresponding to each data-point (X.shape[0])
        """
        if centroids=='kmeans++':
            centroids = initialize_centroids_plus(X,k)
        else:
            centroids = initialize_centroids(X,k)

        # for values of k>20, it takes a very long time to converge
        # so we are putting a limit on iterations such that if k increases,
        # number of iterations decrease
        # also we can have more iterations for smaller values of k, which will
        # yield a better set of centroids
        iter_limit = int(800/k)
        iters = 0
        d = 5
        while (d > tolerance) and (iters< iter_limit):
            bucket_idx = assign_buckets(X, centroids)
            new_centroids = [np.mean(X[bucket_idx == i], keepdims = True, axis = 0) for i in range(k)]
            new_centroids = np.array(new_centroids).ravel().reshape(k,X.shape[1])
            old_centroids = centroids
            centroids = new_centroids
            d = distance_bw_centroids(old_centroids, new_centroids)
            iters+=1
        return centroids,bucket_idx
```

## TOY Dataset with Random Initialization

```

In [17]: #creating dataset
centers = [(-5, -5), (5, 5), (-5,5),(5,-5), (0,0)]
cluster_std = [0.8, 1, 1.5,1, 1]

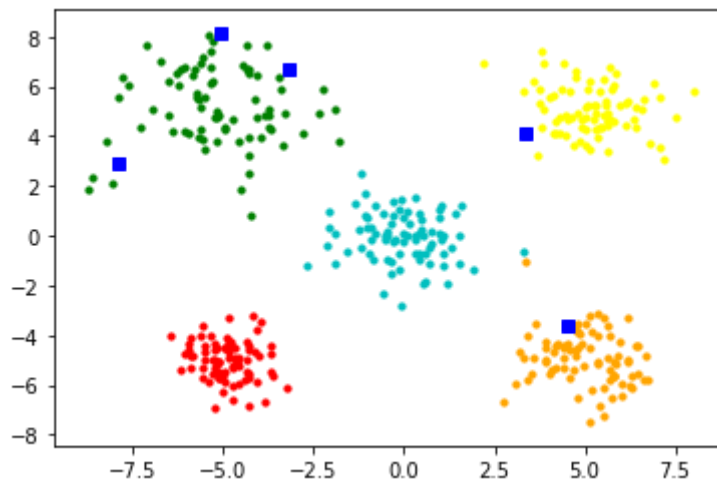
X, y = make_blobs(n_samples=400, cluster_std=cluster_std, centers=centers, n_features=4, random_state=1)

#initial centroids
centroids = initialize_centroids(X,5)

plt.scatter(X[y == 0, 0], X[y == 0, 1], color="red", s=10, label="Cluster1")
plt.scatter(X[y == 1, 0], X[y == 1, 1], color="yellow", s=10, label="Cluster2")
plt.scatter(X[y == 2, 0], X[y == 2, 1], color="green", s=10, label="Cluster3")
plt.scatter(X[y == 3, 0], X[y == 3, 1], color="orange", s=10, label="Cluster4")
plt.scatter(X[y == 4, 0], X[y == 4, 1], color="c", s=10, label="Cluster5")
plt.plot(centroids[:,0],centroids[:,1],"bs")
plt.plot()

```

Out[17]: []

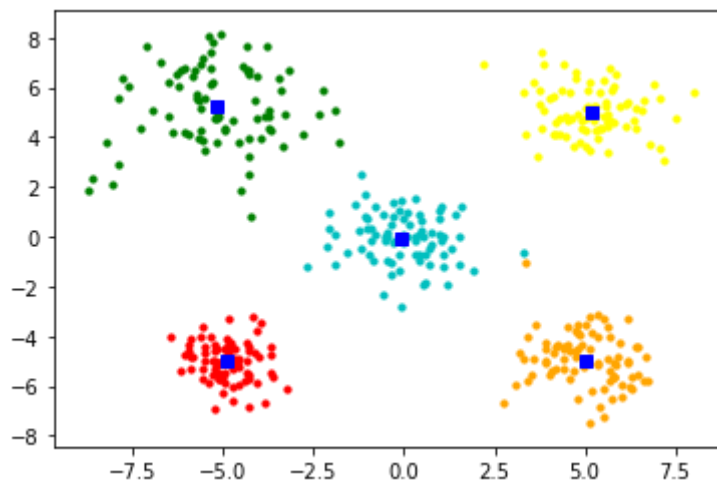


**After applying kmeans with Random Initialization**

```
In [18]: center_kmeans, buckets_kmeans = kmeans(X,k=5)

plt.scatter(X[y == 0, 0], X[y == 0, 1], color="red", s=10, label="Cluster1")
plt.scatter(X[y == 1, 0], X[y == 1, 1], color="yellow", s=10, label="Cluster2")
plt.scatter(X[y == 2, 0], X[y == 2, 1], color="green", s=10, label="Cluster3")
plt.scatter(X[y == 3, 0], X[y == 3, 1], color="orange", s=10, label="Cluster4")
plt.scatter(X[y == 4, 0], X[y == 4, 1], color="c", s=10, label="Cluster5")
plt.plot(center_kmeans[:,0],center_kmeans[:,1], "bs")
```

Out[18]: [



## KMEANS++

As you can see when we randomly initialize, points chosen as initial centroids are far from optimal solution. For larger datasets this may result in slower convergence towards the optimal centroids. As a solution to this problem we use "kmeans++"

How does Kmeans++ initialize centroids?

- 1) Pick a random point from the dataset as the first centroid.
- 2) Find a point in the dataset which is farthest from this point and make it the second centroid.
- 3) Now for selecting the third centroid, for each point in the dataset find the minimum distance of that point from the current set of centroids and then select the point for which this "minimum" distance is the maximum(example below)
- 4) repeat step 3 till we obtain k centroids

Example: Our dataset has 10 points  $p_1, \dots, p_{10}$  we already selected two centroids,  $c_1$  (lets say  $p_3$ ) and  $c_2$  (lets say  $p_7$ ) using the above steps Next we are trying to find the third centroid

MAX (  $\min(p_1-c_1, p_1-c_2), \min(p_2-c_1, p_2-c_2), \min(p_3-c_1, p_3-c_2), \min(p_4-c_1, p_4-c_2), \min(p_5-c_1, p_5-c_2), \min(p_6-c_1, p_6-c_2), \min(p_7-c_1, p_7-c_2), \min(p_8-c_1, p_8-c_2), \min(p_9-c_1, p_9-c_2), \min(p_{10}-c_1, p_{10}-c_2)$  )

Lets assume we get  $p_4-c_2$  after this operation, then  $p_4$  becomes our third centroid

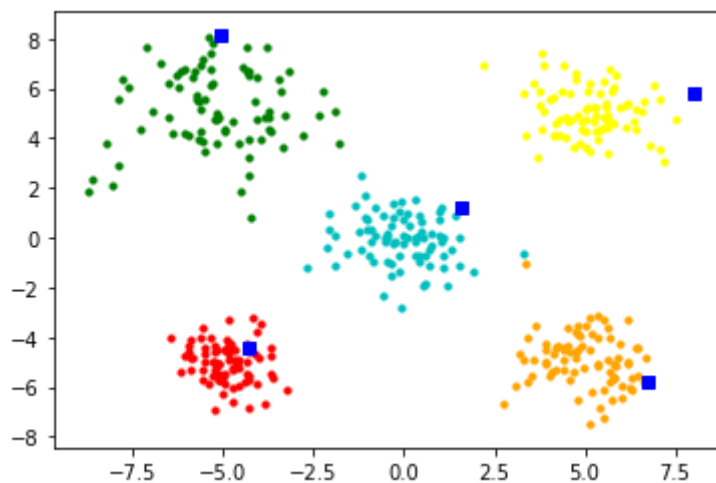
```
In [8]: def initialize_centroids_plus(X:np.ndarray,k:int):
        """
        implementing kmeans++ initialization
        input : X,k
        output : returns k smartly selected centroids
        """
        centroids = initialize_centroids(X,1)
        for _ in range(k-1):
            dist_matrix = np.stack([np.linalg.norm(X-c, axis=1) for c in centroids]) #centroids*X.shape[0]
            min_point_to_centroids = np.min(dist_matrix, axis=0)
            new_centroid_idx = np.argmax(min_point_to_centroids)
            centroids = np.vstack((centroids,X[new_centroid_idx]))
        return centroids
```

## TOY dataset the kmeans++ initialization

```
In [9]: centroids = initialize_centroids_plus(X,5)

plt.scatter(X[y == 0, 0], X[y == 0, 1], color="red", s=10, label="Cluster1")
plt.scatter(X[y == 1, 0], X[y == 1, 1], color="yellow", s=10, label="Cluster2")
plt.scatter(X[y == 2, 0], X[y == 2, 1], color="green", s=10, label="Cluster3")
plt.scatter(X[y == 3, 0], X[y == 3, 1], color="orange", s=10, label="Cluster4")
plt.scatter(X[y == 4, 0], X[y == 4, 1], color="c", s=10, label="Cluster5")
plt.plot(centroids[:,0],centroids[:,1],"bs")
plt.plot()
```

Out[9]: []



Comparing kmeans++ initialization to random initialization, we clearly see that initial locations of the centroids are much closer to their soon to be optimal positions

## MODEL EVALUATION

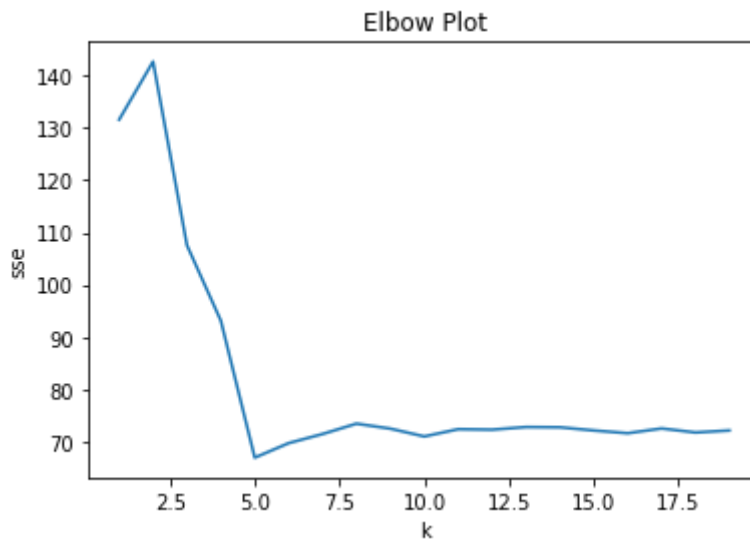
### Elbow Plot

Elbow plot helps us estimate the value of  $k$  which will best fit the data at hand. We choose the value of  $k$  for which either  $sse$  is smallest or after which the plot begins to flatten. The intuition is, we want the  $k$  for which the error is smallest (lowest point in the graph) or we could go for  $k$  beyond which the plot begins to flatten, meaning even if we increase the clusters beyond that point, error is not going to decrease. As discussed above, if we go for more clusters this algorithm starts to get more expensive.

```
In [10]: def sse(X, k):  
         center_kmeans, buckets_kmeans = kmeans(X,k,centroids='kmeans++')  
         sum_err = 0  
         for i in range(k):  
             sum_err += np.linalg.norm(X[buckets_kmeans==i]-center_kmeans[i])  
         return sum_err
```

```
In [11]: x_axis, y_axis = [], []  
         for k in range(1,20):  
             x_axis.append(k)  
             y_axis.append(sse(X,k))  
  
         plt.plot(x_axis, y_axis)  
         plt.xlabel("k")  
         plt.ylabel("sse")  
         plt.title("Elbow Plot")
```

```
Out[11]: Text(0.5, 1.0, 'Elbow Plot')
```



We choose  $k = 5$  in this case

## IMAGE COMPRESSION

One of the interesting applications of clustering is image compression. We take an image which has all sorts of colors and compress to an image with only a few. How do we go about this?

(Please refer to the ipynb in case pdf is not clear) Intuition: 1) We split the image into three channels (r,g and b) 2) For the below 4X3 coloured image, each pixel has r,g and b intensity values. 12 44 12 12 11 14 22 18 23 56 89 91 22 32 01 02 12 72 92 31 81 92 74 65 42 31 15 29 14 23 46 69 22 33 44 81 3) Next we represent each pixel as a vector (r,g,b) In this case P11 -> (12,11,23) P32 -> (31,23,33) and so on 4) Here we will end up with (3\*4, 3) matrix. 5) Kmeans clustering on this dataset will give us a set of centroids. Color associated with each centroid will be the color by which the respective cluster will be identified.

```
In [12]: def compressor(pic:np.ndarray,k:int):
         """
         inputs :
             pic : original image matrix of the form (h,w,c)
             k   : number of clusters
         output : compressed image
         """
         if len(pic.shape)==2: #greyscale images
             h,w = pic.shape
             c =1
         else:
             h,w,c = pic.shape
         pic = pic.reshape(h*w,c)
         centroids, buckets_idx = kmeans(pic,k=k, centroids='kmeans++')
         new_img_arr = centroids[buckets_idx].astype('uint8')
         if c==1:
             new_img_matrix = new_img_arr.reshape(h,w)
         else:
             new_img_matrix = new_img_arr.reshape(h,w,c)
         return Image.fromarray(new_img_matrix)
```



```
In [13]: #image import
im = Image.open('messi.jpg')
pic_array = np.asarray(im)
h,w,c = pic_array.shape
uniq_colors = len(np.unique(pic_array.reshape(h*w,c)))
print(f"number of unique colours in this picture is {uniq_colors}")
im
```

number of unique colours in this picture is 256

Out[13]:



```
In [14]: #compressed image
compressed_pic = compressor(pic_array,20)
pic_array = np.asarray(compressed_pic)
h,w,c = pic_array.shape
uniq_colors = len(np.unique(pic_array.reshape(h*w,c)))
print(f"number of unique colours in this picture is {uniq_colors}")
compressed_pic
```

number of unique colours in this picture is 48

Out[14]:



KIAN

```
In [15]: #grey scale images
im = Image.open('north-africa-1940s-grey.png')
pic_array = np.asarray(im)
im
```

Out[15]:



```
In [16]: compressed_pic = compressor(pic_array,4)
         compressed_pic
```

Out[16]:



Other applications of kmeans clustering could be in extension to a variety of deep learning algorithms. Once we have the embeddings we can treat them as vectors in n-dimensional space and apply kmeans to get similarity between the data points. This acts as a very important tool to analyse a dataset that we know nothing about.